

# Synchronization

## Distributed and Pervasive Systems, MSc

Christian Fischer Pedersen  
cfp@ece.au.dk

Department of Electrical and Computer Engineering  
Aarhus University

Revised on February 5, 2022

# Outline

## **Clock synchronization**

### **Logical clock: Lamport time-stamps**

### **Logical clock: Vector clock**

### **Real clock: Precision Time Protocol**

## **References**

# Outline

## Clock synchronization

Logical clock: Lamport time-stamps


Logical clock: Vector clock

Real clock: Precision Time Protocol

References

# Clock synchronization objectives

- ▶ **Common** notion of local or global, real or logical, **time**
- ▶ Coordination of **logical** or **real** clocks in a distributed system
- ▶ **Logical and real** clocks: Able to obtain logical event ordering
- ▶ **Real clocks**: Ability to meet wall-clock deadlines
- ▶ **Logical clocks**: Ability to meet deadlines relative to events



# Clock synchronization uses

- ▶ Stock market buy and sell orders
- ▶ Secure document timestamps
- ▶ Network gaming
- ▶ Aviation traffic control and position reporting
- ▶ Multimedia synchronization, e.g. real-time teleconf.
- ▶ Event synchronization and ordering
- ▶ Network monitoring, measurement and control
- ▶ ...

## Synchronization

How can you implement a synchronous algorithm in an asynchronous environment?

# Clock synchronization challenges

## Logical clocks

- ▶ Achieve chronological local and global event ordering in a distributed system without synchronous real clocks
- ▶ Initial and on-going corrections are required

## Real clocks

- ▶ Precise synchronization of clocks in nodes
- ▶ Each node's clock tend to drift due to instabilities in source oscillators and environmental conditions such as temperature and mechanical wear and tear
- ▶ Initial and on-going corrections are required

# Clock synchronization solutions for logical clocks

## Examples

- ▶ Lamport Time-Stamps
- ▶ Vector Clocks (AKA vector timestamps or version vectors)
- ▶ Matrix Clocks
- ▶ Plausible Clocks
- ▶ Interval Tree Clocks
- ▶ Bloom Clocks



# Clock synchronization solutions for real clocks

NTP (Network Time Protocol) and SNTP (Simple NTP)

- ▶ Synchronize clocks on the **millisecond** level


RBIS (Reference Broadcast Infrastructure Synchronization)

- ▶ Synchronize clocks on the **microsecond** level


IEEE 1588 Precision Time Protocol (PTP) v1/v2

- ▶ Synchronize clocks on the **nanosecond** level

And many more ...




**Figure:** Local clocks: skew+drift



**Figure:** Global clock: Accurate

## Example: Event ordering: Wrong order



Event Ordering based on Actual Time

Time	Event
3	Event 1
5	Event 2
7	Event 3

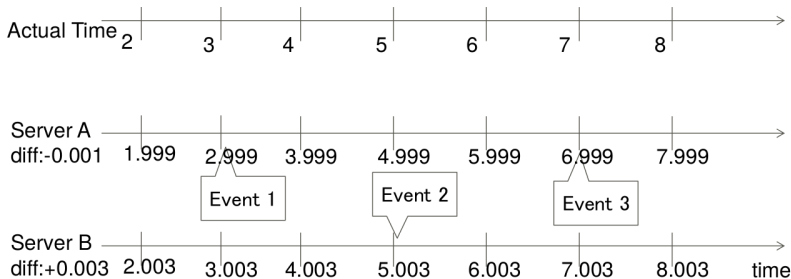
Event Ordering based on Timestamp

Time	Event
2	Event 1
6	Event 3 ←
8	Event 2 ←

reverse!

**Figure:** Reversed events due to offset clocks (Fujitsu).

## Example: Event ordering: Right order



Event Ordering based on Actual Time

Time	Event
3	Event 1
5	Event 2
7	Event 3

Event Ordering based on Timestamp

Time	Event
2.999	Event 1
5.003	Event 2
6.999	Event 3

← correct!

**Figure:** Properly ordered events due to non-offset clocks (Fujitsu).

# Outline

Clock synchronization

**Logical clock: Lamport time-stamps**

Logical clock: Vector clock

Real clock: Precision Time Protocol


References

# Logical clocks in general

- ▶ Achieves chronological event ordering across nodes in a distributed system
- ▶ Especially useful in absence of synchronous real clocks
- ▶ Often it is not necessary to know when (real clock) an event happened, but rather the logical ordering of events
- ▶ Each process keeps track of logical local and global time
- ▶ We need to keep the logical clocks synchronized
- ▶ Typically, a logical clock algorithm updates
  - ▶ logical local time after each local event
  - ▶ logical global time after each data exchange

# About Lamport time-stamps

- ▶ (Leslie) Lamport time-stamps, 1978.  
Turing Award, 2013
- ▶ Algorithm to determine partial ordering of events in a distributed system
- ▶ A Lamport logical clock is basically an incrementing software counter maintained by each process



---

## Time, Clocks, and the Ordering of Events in a Distributed System

Leslie Lamport  
Massachusetts Computer Associates, Inc.

---

## Lamport time-stamps pseudo code

- ▶ A process increments its counter at each local event
- ▶ The counter value is included with the message when sending

### Send

```
time = time+1;  
time_stamp = time;  
send(Message, time_stamp);
```

- ▶ A process updates its counter when receiving a message to the greater of current counter and received time-stamp
- ▶ When the counter incremented, the msg is considered received

### Receive

```
(message, time_stamp) = receive();  
time = max(time_stamp, time)+1;
```

# The Happened-Before Relation 1/2

Events  $\{e_i\}_{i=1}^{E \in \mathbb{Z}_{++}} \subset \mathcal{E}$  are time-stamped by  $C(e_i) \subset \mathbb{Z}_{++}$  such that

1.  $e_1 \rightarrow e_2 \Rightarrow C(e_1) < C(e_2)$  (clock consistency)
2. thus, if  $C(e_1) \not< C(e_2)$  then  $e_1 \not\rightarrow e_2$  (contrapositive)
3.  $e_1 \rightarrow e_2 \Leftrightarrow C(e_1) < C(e_2)$  (strong clock consistency)

where  $\rightarrow$  denotes *happened-before*.

- ▶ Lamport clocks: 1, 2
- ▶ Vector/matrix clocks: 1, 2, 3 (strong clock consistency)

Concurrent events

- ▶  $e_1 \not\rightarrow e_2$  and  $e_2 \not\rightarrow e_1$



# The Happened-Before Relation 2/2

The following holds for the *happened-before* relation ( $\rightarrow$ )

- ▶ Transitive:

$\forall e_1, e_2, e_3$ , if  $e_1 \rightarrow e_2$  and  $e_2 \rightarrow e_3$ , then  $e_1 \rightarrow e_3$

- ▶ Irreflexive:

$\forall e_i : e_i \not\rightarrow e_i$

- ▶ Antisymmetric:


$\forall e_1, e_2$ , where  $e_1 \neq e_2$ , if  $e_1 \rightarrow e_2$ , then  $e_2 \not\rightarrow e_1$ .

## Example: Event counting: The setup

Consider the following

- ▶ Processes:  $P_1, P_2, P_3$
- ▶ Events:  $a, b, c, \dots, k$
- ▶ Local event counter in each process
- ▶ Clock is advanced between consecutive events in same process
- ▶ Processes communicate
- ▶ Each message carries a time-stamp of sender's logical clock

## Example: Event counting: Wrong ordering




Wrong ordering

- ▶  $e \rightarrow h$  but  $5 \geq 2$
- ▶  $f \rightarrow k$  but  $6 \geq 2$

Fig.: Courtesy of P. Krzyzanowski

## Example: Event counting: Corrected ordering



Corrected ordering

- ▶  $e \rightarrow h$  and  $5 < 6$
- ▶  $f \rightarrow k$  and  $6 < 7$

Fig.: Courtesy of P. Krzyzanowski

# Outline

Clock synchronization

Logical clock: Lamport time-stamps

**Logical clock: Vector clock**

Real clock: Precision Time Protocol

References

## Motivation for vector clocks 1/2

Lamport time-stamps induce an order **consistent** with causality

► **but** the converse of the clock condition does **not** hold

Let events  $\{e_i\}_{i=1}^{E \in \mathbb{Z}_{++}} \subset \mathcal{E}$  be time-stamped by  $C(e_i) \subset \mathbb{Z}_{++}$

1.  $e_1 \rightarrow e_2 \Rightarrow C(e_1) < C(e_2)$  (clock consistency)
2. thus, if  $C(e_1) \not< C(e_2)$  then  $e_1 \not\rightarrow e_2$  (contrapositive)
3.  $e_1 \rightarrow e_2 \Leftrightarrow C(e_1) < C(e_2)$  (strong clock consistency)

where  $\rightarrow$  denotes *happened-before*.

- Lamport clocks: 1, 2
- Vector/matrix clocks: 1, 2, 3 (strong clock consistency)

Concurrent events

- $e_1 \not\rightarrow e_2$  and  $e_2 \not\rightarrow e_1$

## Motivation for vector clocks 2/2

If the **strong clock consistency**

▶  $e_1 \rightarrow e_2 \Leftrightarrow C(e_1) < C(e_2)$

does **not** hold, then **it may** be that

▶  $C(e_1) < C(e_2)$  **even if**  $(e_1 \not\rightarrow e_2 \wedge e_2 \not\rightarrow e_1)$  (concur. events)

▶ **But** if events  $e_1$  and  $e_2$  are concurrent, then it **should** be that  $C(e_1) = C(e_2)$

Thus, vector clocks detects causality violations unseen by Lamport timestamps

- ▶ We need a clock mechanism that is **necessary** (Lamport) **and sufficient** (Vector clocks) in capturing causality

## The vector clock algorithm in words

A vector clock of a system of  $N$  processes is a vector of  $N$  logical clocks, one clock per process

- ▶ Initially all clocks are zero.
- ▶ Each time a process experiences an internal event, it increments its own logical clock in the vector by one.
- ▶ Each time a process sends a message, it increments its own logical clock in the vector by one and then sends a copy of its own vector.
- ▶ Each time a process receives a message, it increments its own logical clock in the vector by one and updates each element in its vector by taking the element-wise maximum of the value in its own vector clock and the value in the vector in the received message



# Vector clock example




Fig.: Courtesy of P. Krzyzanowski

# Outline

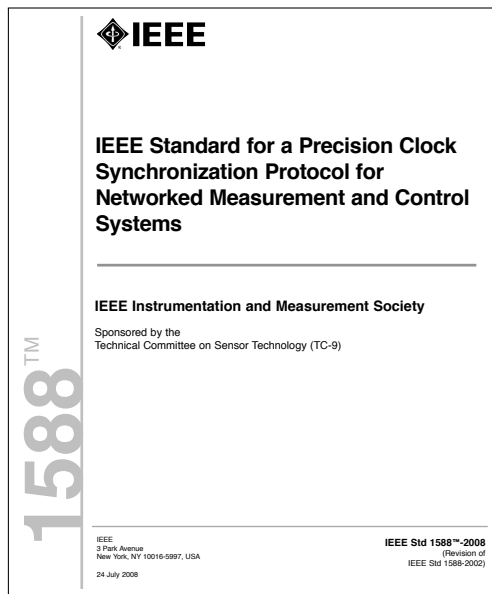
Clock synchronization

Logical clock: Lamport time-stamps

Logical clock: Vector clock

**Real clock: Precision Time Protocol**


References



**Figure:** Accessible via IEEE Xplore (289 pages).

# Clock synchronization by Precision Time Protocol

- ▶ To synchronize clocks in nodes on a computer network
- ▶ Achieves **sub- $\mu$ s** clock synchronization accuracy
- ▶ Suitable for distributed real-time systems
- ▶ Works on packet switched networks, e.g. Ethernet



**Figure:** Clock synchronization by the Precision Time Protocol

# Precision Time Protocol: Timeline

- ▶ Originally defined in the IEEE 1588-2002 standard (PTPv1)
- ▶ Revised in 2008 to IEEE 1588-2008 (PTPv2)
- ▶ PTPv2 improves accuracy, precision, and robustness
- ▶ PTPv2 is **not backwards compatible** with PTPv1
- ▶ Current information may be found at [nist.gov](http://nist.gov) and [ieee.org](http://ieee.org)
- ▶ Conferences on IEEE 1588 PTP held in 2003, 2004, ...

# Precision Time Protocol: E.g. application area

## Industrial automation

- ▶ Plant floor control and management of sensors and actuators

## Telecommunication


- ▶ Control voice, video, and text services in beyond 3G arch.

## Audio and video bridging (IEEE AVB 802.1)

- ▶ Time-sensitive interoperability between multimedia devices

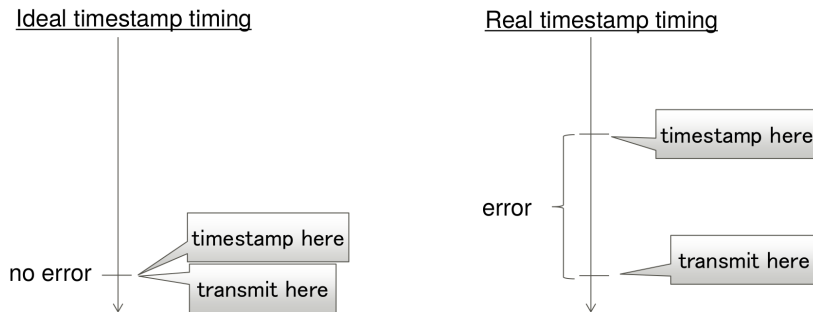
# Precision Time Protocol: Overall mechanics

- ▶ For packet switched networks, e.g. Ethernet
- ▶ Synchronization and data transfer on the **same** network
- ▶ Minimal admin, network, software and hardware requirements
- ▶ Most precise clock synchronizes the other clocks



(Weibel, 2009 [2])

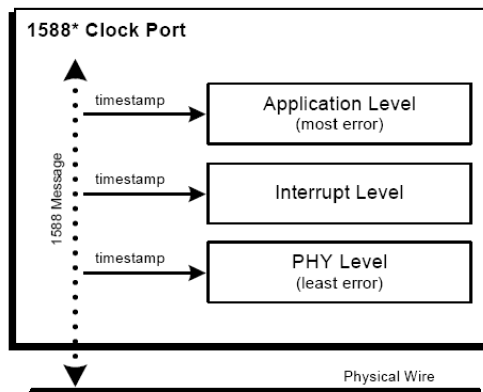
# Time stamping for clock synchronization



**Figure:** Minimize the time (error) between time of stamping and transmitting or receiving (Fujitsu).

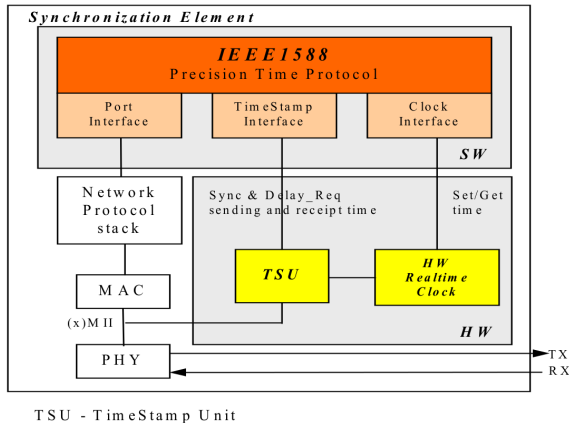


# HW vs. SW time stamping



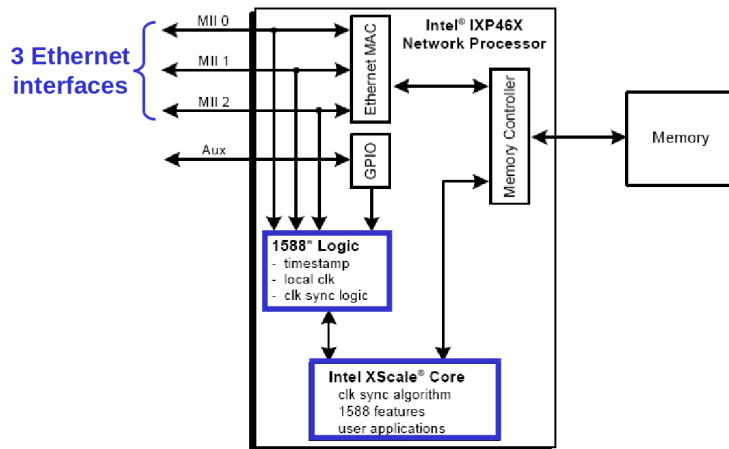
**Figure:** Hardware vs. software time stamping (Intel).

# Synchronization element with Time Stamp Unit




**Figure:** Synchronization element with Time Stamp Unit (TSU) and Real-time Clock in hardware. The rest of the protocol can be in software as timing requirements are lower (Mohl, 2003 [1]).

# Hardware support for IEEE PTP



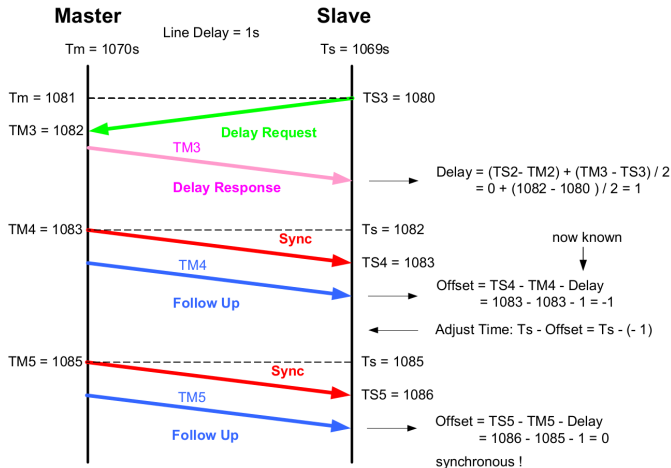
**Figure:** Hardware support for IEEE PTP, e.g. Intel IXP46X Network Processors (Intel).

# PTP: Phase 1: Offset correction



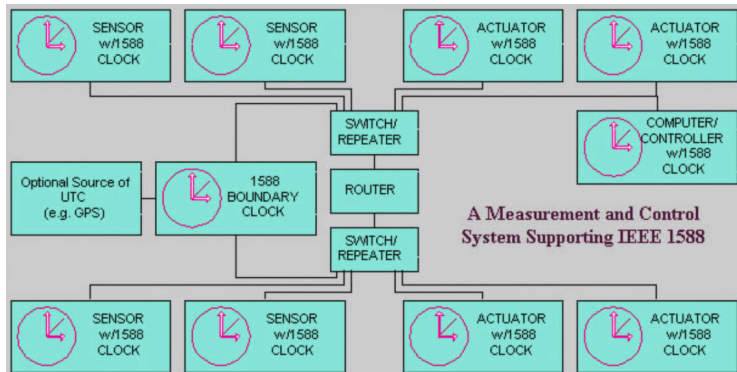
**Figure:** In offset correction, the line delay is assumed to be 0s. The actual delay (1s) is measured and added in the next phase (Mohl, 2003 [1]).

# PTP: Phase 2: Delay correction



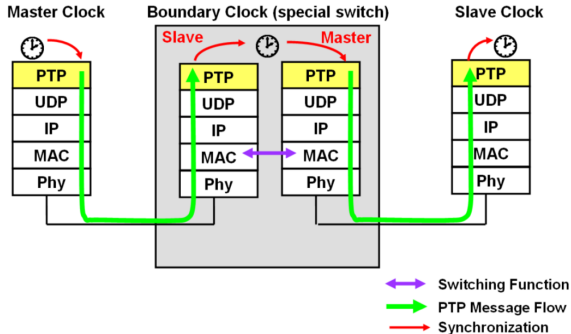
**Figure:** Known: M is  $\geq 0s$  ahead. Assumed: Symmetric delay (Mohl, 2003 [1]).

## E.g: Two subnets and a boundary clock



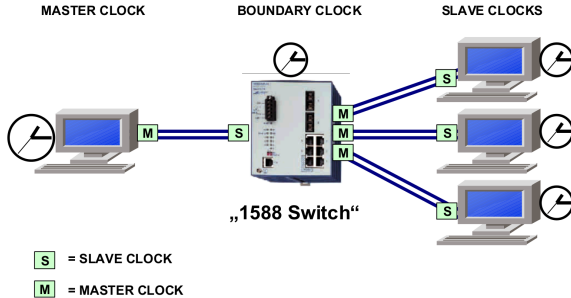
**Figure:** Measurement and control system with IEEE PTPv2 ([www.nist.gov](http://www.nist.gov)).

# Switch with boundary clock 1/2



**Figure:** Boundary clocks synchronize subnet nodes' clocks across switches and routers (Weibel, 2009 [2]).

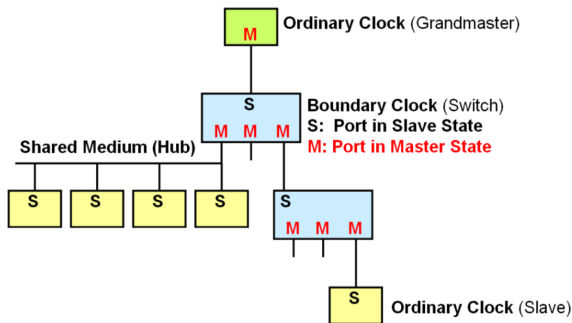
## Switch with boundary clock 2/2



**Figure:** Switch with boundary clock (Mohl, 2003 [1]).




# Synchronization hierarchy



**Figure:** The most precise clock, i.e. the Grandmaster, is selected by the Best Master Clock Algorithm. Grandmaster is root of the hierarchy (Weibel, 2009 [2]).

# Master/Slave synchronization test (PTPv1)


- ▶ Network modules with pulse output (PPS: Pulse Per Second)
- ▶ Oscilloscope measures deviation between master and slave
- ▶ Offset mean, max, and std.dev.:  $(-4.25, \pm 100, 23.95)$  ns



**Figure:** Master/Slave clock setup with Ethernet packet generator

(Mohl, 2003 [1])

# Normal distribution - How errors often distribute



**Figure:** Normal dist: values less than one std,  $\sigma$ , away from the mean,  $\mu$ , account for 68.27% of the set; while  $2\sigma$  from  $\mu$  account for 95.45%; and  $3\sigma$  account for 99.73%. Fig.: Courtesy of Dan Kernler

$$\text{PDF: } f(x|\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

# Outline

Clock synchronization

Logical clock: Lamport time-stamps

Logical clock: Vector clock

Real clock: Precision Time Protocol

**References**

# References I

- [1] Mohl, D. (2003). IEEE 1588 - Precise time synchronization as the basis for real time applications in automation. Technical report, White Paper, Industrial Networking Solutions, pp. 1–8.
- [2] Weibel, H. (2009). Technology update on IEEE 1588: The second edition of the high precision clock synchronization protocol. Technical report, Zurich University of Applied Sciences, Winterthur, Switzerland.